

UDIT SHARMA

AI Engineer · Agentic Systems & LLM Orchestration

Bhopal, India · +91 86020 49101 · uditsharma9981@gmail.com · linkedin.com/in/udit-sharma · github.com/uditstocks

SUMMARY

Agentic AI engineer building production-grade autonomous systems across multi-agent orchestration, RAG pipelines, and LLM workflow design. Open-source contributor to LangChain with merged pull requests to its official documentation, and currently delivering real-world AI systems as an AI Developer at Node and Edges LLP. Focused on translating advanced agentic-AI research into reliable, shippable software.

TECHNICAL SKILLS

Languages	Python, TypeScript · Familiar: JavaScript, React, Rust
AI Frameworks	LangChain, LangGraph, CrewAI, Pydantic AI, Agno, FastAPI, LangServe
AI Architecture	Agentic AI System Design, Multi-Agent Systems, LLM Workflow Orchestration, RAG & Hybrid RAG Pipelines, Tool-Calling Frameworks, Memory Systems, Parallel Agent Execution, Production-Grade AI Application Development
LLM Providers	Anthropic, Google Gemini, Groq, NVIDIA NIM
Tools & Platforms	SQL / PostgreSQL, Vector Databases (pgvector, FAISS), Docker, Kubernetes (basic), Git, GitHub Actions, CI/CD, LangSmith

EXPERIENCE

- AI Developer Intern** · Node and Edges LLP Jun 2026 – Present
- Building and shipping production-grade AI systems end-to-end, delivering real-world projects across the agentic-AI and LLM application stack as part of a professional engineering team.
 - Working within a disciplined engineering workflow - version control, code review, and CI/CD - to ship reliable, maintainable, and observable AI features.

OPEN-SOURCE CONTRIBUTIONS

- LangChain** · Open-Source Contributor - [langchain-ai/docs](https://github.com/langchain-ai/docs)
- Authored a comprehensive comparison of PGVectorStore vs. PGVector for the official documentation, clarifying the architectural trade-offs and helping developers choose the right Postgres-backed vector store on a docs site used by thousands of developers. [PR #3742](#)
 - Documented the previously-undocumented `set_entry_point` and `set_finish_point` lifecycle methods in LangGraph, enabling developers to build graph-based agent workflows correctly without resorting to trial-and-error. [PR #3833](#)
 - Corrected a misleading `delete_old_messages` example in the message-management docs, preventing developers from shipping broken message-retention logic to production. [PR #4009](#)
 - Identified and reported a critical CI outage affecting every contributor to the repository - a deprecated, unmaintained Vale Snap package failing mid-build - and submitted a surgical fix: a version-pinned (v3.4.1) direct binary download replacing the Snap dependency, which passed all CI checks on the first attempt and restored reliable builds for the whole project. [Issue #3837](#) [PR #3841](#)

PROJECTS

Trinetra Capital AI - Autonomous Multi-Agent Trading System

Python · LangGraph · LangChain · NVIDIA NIM · LangSmith · Docker

- Architected a hierarchical multi-agent system on LangGraph's supervisor pattern: a Groq LLaMA-3.3-70B supervisor interprets user intent and routes work to specialized Research and Trading agents (NVIDIA NIM Nemotron-3-Super-120B) through agent-to-agent handoff tools, with StateGraph-driven control flow and checkpointed conversation state.

- Engineered a human-in-the-loop safety layer using LangChain’s HumanInTheLoopMiddleware and LangGraph interrupt(): every simulated buy/sell order is paused for explicit approve/reject review - with full symbol, share count, price, and total surfaced - before it can execute.
- Built real-time, multi-market equity research over Yahoo Finance spanning Indian (NSE/BSE) and US exchanges, with automatic exchange and currency detection (₹ INR / \$ USD) and live retrieval of price, P/E ratio, market cap, 52-week range, and sector.
- Reduced token consumption ~10× (~16K–60K → ~2K–6K per run) through prompt and workflow optimization, instrumented end-to-end with LangSmith for per-tool token analytics, tool-I/O tracing, and P50/P99 latency monitoring.
- Implemented persistent portfolio tracking that survives sessions and container restarts, and packaged the full stack with Docker / Docker Compose for reproducible, one-command deployment.

PageIndex - Vectorless RAG Pipeline

Python · LangChain · Retrieval-Augmented Generation

- Built a “vectorless” RAG pipeline that answers questions over documents with no vector database and no embeddings - replacing semantic vector search with PageIndex-style structural/keyword retrieval (chunk → retrieve → LLM).
- Removed the embedding model and vector store from the stack entirely, eliminating a significant source of infrastructure, cost, and setup complexity while keeping every answer grounded in the source documents.
- Implemented in Python with LangChain as a lightweight, low-dependency, production-aware pipeline that prioritizes fast setup and simple deployment over heavyweight retrieval infrastructure.

EDUCATION

B.Tech, Artificial Intelligence - Technocrats Institute of Technology, Bhopal

Expected 2027

Specialization: AI & Computer Science Engineering · Coursework: Agentic AI, Generative AI, RAG Systems, LLM Workflow Design, Python, NLP, Data Structures & Algorithms

Senior Secondary (Class XII) — Maria Mata Sr. Sec. School, Chhatarpur

2023

CERTIFICATIONS

- **IBM** - Generative AI, Artificial Intelligence Fundamentals
- **Red Hat** - AI Foundations Technologist Certificate
- **LangChain** - Foundations to LangGraph
- **CrewAI** - AI Agents System
- **Cisco** - Python Essentials 1 & Python Essentials 2

ADDITIONAL

Languages

English (Professional), Hindi (Native)

Interests

Agentic AI · Equity research & long-term investing · Reading company financials

Open to

AI / agentic-engineering roles and open-source collaboration