

UDIT SHARMA

AI Engineer (B.Tech AI, 2027) · Agentic Systems, MCP & Retrieval

Bhopal, India · +91 86020 49101 · uditsharma9981@gmail.com · linkedin.com/in/udit-sharma-b89995290 · github.com/uditstocks

SUMMARY

AI Engineer who designs and ships agentic systems, MCP servers and retrieval pipelines in Python. Creator of Trinetra Capital AI, a multi-agent trading system Claude and ChatGPT drive directly through MCP. Open-source contributor with five merged PRs to LangChain's official documentation. Builds production-grade autonomous systems across multi-agent orchestration, RAG pipelines and LLM workflow design.

TECHNICAL SKILLS

Languages	Python (primary), TypeScript, SQL · familiar: JavaScript, React, Rust
Agents & MCP	LangGraph multi-agent orchestration (supervisor, sub-agents, handoffs, HITL interrupts, checkpointing) · LangChain · LlamaIndex · MCP server development (Model Context Protocol, FastMCP, tools / resources / prompts, stdio & Streamable HTTP, OAuth 2.1 remote MCP) · tool & function calling · AI assistants & copilots · Claude Agent SDK & Agent Skills (exploring)
LLM Engineering	context engineering · structured outputs · LLM evaluation · tracing & observability (LangSmith, Langfuse) · guardrails & prompt-injection defense · Responsible AI · fine-tuning · model routing for cost & latency · OpenAI-compatible APIs · async Python (asyncio, FastAPI)
Retrieval Models	Retrieval-Augmented Generation (RAG), hybrid search, reranking, chunking, vector databases (pgvector, FAISS) · Anthropic Claude, Google Gemini, NVIDIA NIM (Nemotron, Llama), Groq, OpenRouter, Ollama, Hugging Face · exploring: OpenClaw
Backend & Infra	FastAPI (REST APIs), PostgreSQL, SQLAlchemy, Alembic, OAuth 2.1 / JWT, Docker, Kubernetes (basic), Railway deployment, CI/CD (GitHub Actions), unit and integration testing (pytest)

EXPERIENCE

AI Developer Intern · Node and Edges LLP Jun 2026 – Present

- Design multi-agent architectures and develop AI features end to end, from requirements and implementation through tests, documentation and code review; shipped two reviewed and merged pull requests to [mcplint](#), the company's MCP security linter.
- Provide high-level AI advisory and consultancy to the team and stakeholders: evaluating frameworks and tooling, recommending architectures, and translating business requirements into technical plans and delivery timelines.

Developer Relations Engineer, Forge-Guild Cipher · P2P.me · part-time, remote Jun 2026 – Present

- Contribute code to P2P.me's open-source p2pkid and payment-integrator repositories, alongside technical documentation, reproducible bug reports and localization; weekly written progress reports, deliverables kept on schedule.
- Take part in international team and client meetings, owning client communication and follow-ups across time zones in a distributed team.

PROJECTS

Trinetra Capital AI · MCP server: equity research and human-approved trading on NSE/BSE github.com/uditstocks/trinetra-capital-ai-mcp

- 19 tools, 3 resources, 3 prompts; the host model reasons while the server stays deterministic, with no model call or model key, and returns indicator-level reasoning traces. Replaced an earlier LangGraph supervisor CLI once the host model made it redundant.
- Two-step orders (preview, then a single-use expiring token; one tool trades); caps, daily limits, kill switch and trading mode enforced in the service layer, never from tool arguments, so no host call bypasses them; hash-chained audit log; AES-256-GCM encrypted credential vault.

VectorlessRAG · RAG app built on VectifyAI's open-source PageIndex indexer, no vector database github.com/uditstocks/PageIndex

- Built the four-stage Librarian, Navigator, Reader, Generator query pipeline, with human confirm/override/skip at every stage, on top of PageIndex's TOC-tree indexer (upstream indexing kept intact); routed all model calls through LiteLLM to NVIDIA NIM.
- LLM page selection with stated reasoning, page-level citations or explicit refusal, robust parsing of malformed model output; evaluation docs state what is measured (indexing accuracy via upstream's `verify_toc` gate) and what is not yet (recall@k, faithfulness, latency).

OPEN-SOURCE CONTRIBUTIONS

LangChain · official documentation, langchain-ai/docs · 5 merged PRs github.com/langchain-ai/docs/pulls?q=author:uditstocks

- Diagnosed a CI outage that was intermittently failing checks on every contributor's pull request, traced it to a deprecated Vale snap install, and shipped a version-pinned binary fix that passed all CI checks; maintainers landed the identical fix the same day #3837 · #3841.
- Fixed two tutorials that broke as written: the LangGraph SQL-agent guide, where a loop variable shadowed the `@tool` decorator and raised `TypeError` at the human-in-the-loop step #4727, and the Text Embeddings Inference guide, rewritten to the OpenAI-compatible endpoint after langchain-huggingface began rejecting URLs #4726.
- Three more merged PRs across LangGraph and vector-store documentation (`set_entry_point` / `set_finish_point`, `delete_old_messages`, PGVectorStore vs PGVector), plus two RAG how-to guides open for review #4723 · #4724.

EDUCATION

B.Tech, Artificial Intelligence Expected 2027
Technocrats Institute of Technology, Bhopal
Coursework: Agentic AI, Generative AI, RAG Systems, NLP, DSA
Languages English (professional), Hindi (native)

CERTIFICATIONS

- IBM** · Generative AI; Artificial Intelligence Fundamentals
- Red Hat** · AI Foundations Technologist
- LangChain** · Foundations to LangGraph
- CrewAI** · AI Agents System · **Cisco** · Python Essentials 1 & 2